


Queueing latency (lower is better)



Tail per-token cost (lower is better)

